

# Niko Strongioglou

☎ +41(0)787042239

📍 Salerstrasse 16, 8050 Zurich

🌐 [web](#) [in linkedIn](#) [github](#) [leetCode](#) [email](#)

I am a software engineer with 10 years of experience building large-scale, data-intensive, distributed systems. I revel in being a member of teams composed of smart, self-motivated, and inspired people who consider Computer Science as a means of changing the world.

## Professional Experience

### Senior Software Engineer at 1PlusX/TripleLift *Zurich, Switzerland* May 2020 - Present

Data management and analytics platform providing multi-dimensional batch and real-time user tracking and profiling. Data consolidation and Audience teams.

- Management of various multi-sprint/multi-quarter and cross-team data projects: Leading and mentoring small groups of developers while resolving ambiguity with non-engineering dependencies (e.g. legal, product, customer success, external clients)
- Created or maintained 50+ unique data pipelines running on production systems with more than 2PB of daily loads (spawning 30-40k Spark jobs daily).
- Extensive work on AWS. Developed and maintained Terraform infra of hundreds of services.
- Tag manager ingestion restructuring resulting in a data-set specific size decrease of up to 80% - while maintaining targeting quality.
- Decreased pipeline runtime up to 15% by introducing batch REST API endpoints.
- Close collaboration with ML engineers to deliver useful data for their pipelines - MLOps.
- Cross-team user ID storage optimization resulting in up to 20% decrease in Couchbase loads. Load on external APIs was decreased up to 40%.
- Data privacy pipelines: Access and deletion requests' handling of multiple protocols (GDPR, CCPA).
- DSL design for ETL ops: Rich syntax for creating and combining attributes.
- Real-time data ingestion infrastructure design able to accommodate 7.5M unique events per second.
- E2E user tracking: From cookies and browser-side SDKs to graph-based ID management.
- UI work on analytics dashboards, widgets, data pipeline builder.

Scala Java Go Python Typescript Angular 2+ Redux Akka Spark Luigi AWS EMR ECS  
EC2 ECR EKS RDS S3 CloudWatch Kafka Docker Terraform Packer Grafana InfluxDB  
Couchbase DynamoDB PostgreSQL Databricks Elasticsearch Opensearch

### Software Engineer at Veezoo AG *Zurich, Switzerland* March 2019 - May 2020

NLP platform providing analytics through a variety of corporate data source integrations - namely Oracle, MySQL, PostgreSQL, Salesforce, IBM DB2, Hadoop, SAP HANA - using a QA interface.

- Knowledge graph ontology representations API for (non)-relational data sources which automated data onboarding.
- Semantic parsing pipelines for natural languages (bag of words to logical forms to SQL generation).
- Extend and maintain the [Duckling](#) language engine

Scala Play Apache Jena MySQL PostgreSQL Solr Haskell Docker OpenShift Duckling

### Software Engineer at Locatee AG *Zurich, Switzerland* April 2017 - March 2019

Analytics platform providing hyperlocation, real-time workspace occupancy insights and predictive APIs through geospatial data scraping (SNMP protocol).

- Led the modularization of the core app monolith. From creating three distinct services - Analytics API, Data processing service, and pollers - to POCs and eventually to global roll-outs.
- Online data distillation pipeline detecting outliers on workspace behavior and identifying team collaboration coefficients.
- Different MySQL read/write optimizations for measurements post-processing

Java 8 Spring Core MySQL InfluxDB Tableau Neo4J Kafka Python Jupyter Pandas Docker

## Software Engineer at Upstream Systems SA *Athens, Greece*      October 2014 - March 2017

Green-field project re-engineering the core CRM platform as a fully distributed system. The implementation included a multitude of APIs for data streaming and job scheduling accommodating thousands to millions of transactions daily.

- Asynchronous event-driven services on distributed environments using Guava and RabbitMQ. High availability and resiliency optimizations.
- Implementation of Elasticsearch user samplers via Groovy scripts and custom treemap algorithms. Stretch testing, benchmarking, and performance tweaking on Elasticsearch operations.
- Performance profiling (Dropwizard's Metrics, YourKit, Grafana, and other tools).
- API versioning (URL & Resource Representation).
- Batch job framework implementation with RabbitMQ and Quartz. Throttling and flow control features included.
- QA Automation framework (Selenium, Spring).
- [Recommendation engine](#) using Mahout and supervised learning regression algorithms (Log Likelihood Similarity for item-based recommender and Pearson correlation similarity for user-based recommender).

Java 8   Spring Core   Spring Boot   Spring JPA   Spark   TestNG   RabbitMQ   PostgreSQL   EhCache  
Guava   Quartz   Flyway   Hibernate   Elasticsearch   Camel   Kafka   Spark   Redis   Docker   Ansible  
Protobuf

## Software Engineer at Accenture *Athens, Greece*      June 2014 - September 2014

Provisioning and mediation over Oracle databases and mobile telecom software.

TOAD   SQL Developer   WebLogic

## Freelancing Web Developer *Athens, Greece*      January 2012 - November 2013

Software development on static content pages and web applications.

HTML / CSS   Bootstrap   Javascript   jQuery   Angular JS

## Other Works

---

- Designer and developer of a [B+ indexed RDBMS in C](#).
- Co-creator of a [social network](#) based on node.js and socket.io. The network provided real-time chatting and financial transactions between users.
- Various API prototypes using Flask and FastAPI. [\[1\]](#) [\[2\]](#)

## Education

---

### Degrees

September 2013	<b>M.Sc. in Computer Science</b> Department of Computer Science, National and Kapodistrian University of Athens - 7.49/10 or Upper Second-Class Honours
----------------	--

### Languages

Greek	Native
English	Certificate of Proficiency in English, University of Cambridge, Grade A Certificate of Proficiency in English, University of Michigan, Honors
German	B1
French	Diplome d'études en langue Francaise (DELF)

## References

---

Available upon request