

Niko Strongioglou

☎ +41(0)787042239

📍 Salerstrasse 16, 8050 Zurich

🐙 [github](#) [in linkedIn](#) 🌐 [web](#) 📁 [leetCode](#) ✉ [email](#) 📝 [blog](#)

Technical leader with 10+ years of experience designing large-scale distributed systems and data platforms at global scale. Working across disciplines to deliver strategic impact.

Professional Experience

Lead Software Engineer at ethonai Zurich, Switzerland January 2025 - Present
Platform team lead (6 engineers). Owning MLOps and data, serving vision, time-series anomaly detection, and causal inference teams.

- Redesigned company-wide data integration: Introduced a scalable, EKS-based, and configurable worker paradigm that decreased the effort of FDEs from days to hours and expanded the product to new markets.
- Established platform standards for raw data parsing and loading pipelines, allowing composable pre-/post-processing and inference. Support for 100+ data and metadata tables.
- Architected infra for fine-tuning vision foundation models using PyTorch and AWS Batch. ([DINO](#), [SAM 2](#))
- Designed and developed vision inference pipelines with FastAPI, k8s and Triton (point-to-region, region type identification, anomaly detection/classification) - 2ms runtimes.
- Provided technical leadership and oversight for the migration from legacy single-node docker-compose deployments to Amazon EKS.
- Data management optimizations: compression (up to 90%+), query optimization for training data retrieval (minutes to seconds), continuous aggregates using materialized views, [CloudNativePG](#) using [pgpool-II](#) routing, passive/active and read/write nodes, customer data segmentation.
- Defined engineering & hiring standards by conducting 40+ technical interviews.

Senior Software Engineer at ethonai Zurich, Switzerland June 2024 - December 2024
MLOps and data on computer vision pipelines. Member of the visual quality inspection team.

- Designed and led the implementation of a distributed system for exchanging vision models and anomaly detection results between training servers and edge devices.
- Redesigned vision data storage architecture and reduced storage footprint by 40%.
- Observability standards for vision pipelines, pertaining to ROI scores, defect heatmaps, reference images, annotations, model evaluation, other salient ([SIFT](#)) features.
- Drove improvements to testing infrastructure: UI unit testing, revised integration test framework.

Python Go PostgreSQL Timescale FastAPI Flask [Triton](#) Angular TensorFlow ONNX Hydra
CUDA MQTT Docker k8s [kind](#) [DevSpace](#) [Aim](#) [Gradio](#) [MetaFlow](#) [DVC](#) PyTorch Buildkite
[ArgoCD](#) AWS GCP Azure

Senior Software Engineer at 1PlusX/TripleLift Zurich, Switzerland May 2020 - April 2024
Data management and analytics platform providing multi-dimensional batch and real-time user tracking and profiling. Data consolidation and Audience teams.

- Led multi-sprint/multi-quarter and cross-team data projects: Mentoring small groups of developers while resolving ambiguity with non-engineering dependencies such as legal, product, customer success and external clients.
- Created or maintained 50+ unique data pipelines running on production systems with more than 2PB of daily loads (spawning 30-40k Spark jobs daily).
- Developed and maintained Terraform infra for hundreds of AWS services.
- Tag manager ingestion restructuring resulting in a data-set specific size decrease of up to 80% - while maintaining targeting quality.
- Decreased pipeline runtime up to 15% by introducing batch REST API endpoints.
- Defined contracts between data engineering and ML teams, reducing iteration cycles and schema drifts.
- Cross-team user ID storage optimization resulting in up to 20% decrease in Couchbase loads. Load on external APIs was decreased up to 40%.
- Data privacy pipelines: Access and deletion requests' handling of multiple protocols (GDPR, CCPA).
- DSL design for ETL ops: Rich syntax for creating and combining attributes.
- Real-time data ingestion infrastructure design able to accommodate 7.5M unique events per second.

- E2E user tracking: From cookies and browser-side SDKs to graph-based ID management.
- UI work on analytics dashboards, widgets, data pipeline builder.

Scala Java Go Python Typescript Angular 2+ Redux Akka Spark Luigi AWS EMR ECS
EC2 ECR EKS RDS s3 CloudWatch Kafka Docker Terraform Packer Grafana InfluxDB
Couchbase DynamoDB PostgreSQL Databricks Elasticsearch OpenSearch

Software Engineer at Veezoo AG Zurich, Switzerland

March 2019 - May 2020

NLP platform providing analytics through a variety of corporate data source integrations - namely Oracle, MySQL, PostgreSQL, Salesforce, IBM DB2, Hadoop, SAP HANA - using a QA interface.

- Knowledge graph ontology representations API for (non)-relational data sources which automated data onboarding.
- Semantic parsing pipelines for natural languages (bag of words to logical forms to SQL generation).
- Extend and maintain the Duckling language engine.

Scala Play Apache Jena MySQL PostgreSQL Solr Haskell Docker OpenShift Duckling

Software Engineer at Locatee AG Zurich, Switzerland

April 2017 - March 2019

Analytics platform providing hyper-location, real-time workspace occupancy insights and, predictive APIs through geospatial data scraping (SNMP protocol).

- Led the modularization of the core app monolith. From creating three distinct services - Analytics API, Data processing service, and pollers - to POCs and eventually to global roll-outs.
- Online data distillation pipeline detecting outliers on workspace behavior and identifying team collaboration coefficients.
- Different MySQL read/write optimizations for measurements post-processing

Java 8 Spring Core MySQL InfluxDB Tableau Neo4J Kafka Python Jupyter Pandas Docker

Software Engineer at Upstream Systems SA Athens, Greece

October 2014 - March 2017

Green-field project re-engineering the core CRM monolith as a fully distributed system. The implementation included a multitude of APIs for data streaming and job scheduling accommodating thousands to millions of transactions daily.

- Asynchronous event-driven services on distributed environments using Guava and RabbitMQ. High availability and resiliency optimizations.
- Implementation of Elasticsearch user samplers via Groovy scripts and custom treemap algorithms. Stretch testing, benchmarking, and performance tweaking on Elasticsearch operations.
- Performance profiling (Dropwizard's Metrics, YourKit, Grafana, and other tools).
- API versioning (URL & Resource Representation).
- Batch job framework implementation with RabbitMQ and Quartz. Throttling and flow control features included.
- QA Automation framework (Selenium, Spring).
- [Recommendation engine](#) using Mahout and supervised learning regression algorithms (Log Likelihood Similarity for item-based recommender and Pearson correlation similarity for user-based recommender).

Java 8 Spring Core Spring Boot Spring JPA Spark TestNG RabbitMQ PostgreSQL EhCache
Guava Quartz flyway Hibernate Elasticsearch Camel Kafka Spark Redis Docker Ansible
Protobuf

Software Engineer at Accenture Athens, Greece

June 2014 - September 2014

Provisioning and mediation over Oracle databases and mobile telecom software.

TOAD SQL Developer WebLogic

Freelancing Web Developer Athens, Greece

January 2012 - November 2013

Software development on static content pages and web applications.

HTML / CSS Bootstrap JavaScript jQuery AngularJS

Other Works

- Designer and developer of a [B+ indexed RDBMS in C](#).
- Co-creator of a [social network](#) based on node.js and socket.io. The network provided real-time chatting and financial transactions between users.
- Prototypes for content moderation and text search using Flask and FastAPI. [\[1\]](#) [\[2\]](#)

Education

Degrees

September 2013	M.Sc. in Computer Science Department of Computer Science, National and Kapodistrian University of Athens - 7.49/10 or Upper Second-Class Honours
----------------	--

Languages

Greek	Native
English	Certificate of Proficiency in English, University of Cambridge, Grade A Certificate of Proficiency in English, University of Michigan, Honors
German	B1
French	Diplome d'etudes en langue Francaise (DELF)

Conferences

- [MLOps World](#) 2025

References

Available upon request