

Niko Strongioglou

☎ +41(0)787042239

📍 Salerstrasse 16, 8050 Zurich

🐙 [github](#) [in linkedIn](#) 🌐 [web](#) 🏠 [leetCode](#) ✉ [email](#) 📝 [blog](#)

Technical leader with 10+ years of experience designing large-scale distributed systems and data platforms at global scale. Working across disciplines to deliver strategic impact.

Professional Experience

Lead Software Engineer at ethon Zurich, Switzerland

November 2024 - November 2026

Platform team lead (6 engineers). Owning MLOps and data, serving vision, time-series anomaly detection, and causal inference teams.

- Redesigned company-wide data integration: Introduced a scalable, EKS-based, and configurable worker paradigm that decreased the effort of FDEs from days to hours and expanded the product to new markets.
- Established platform standards for raw data parsing and loading pipelines, allowing composable pre-/post-processing and inference. Support for 100+ data and metadata tables.
- Architected infrastructure for fine-tuning vision foundation models using PyTorch and AWS Batch. ([DINO](#), [SAM 2](#))
- Introduced dedicated, MIG, and time-sliced GPU pools - spanning multiple AZs - with [resource flavor](#) fallback to address GPU scarcity for Kubernetes training workloads.
- Designed and led implementation of vision inference pipelines with FastAPI, Kubernetes, and [Triton](#) (point-to-region, region type identification, anomaly detection/classification) - 2ms runtimes.
- Provided technical leadership and oversight for the migration from legacy single-node docker-compose deployments to EKS, defining deployment, scaling, security and operational patterns adopted across teams.
- Guided the architecture of a distributed ML experimentation and evaluation platform (hundreds of datasets) using [DVC](#), AWS Batch (GPU) with experiment tracking via [mlflow](#).
- Drove database optimizations including compression (up to 90%+ or $\approx$ \$200 savings per customer), query runtimes for training data retrieval (minutes to seconds), continuous aggregates using materialized views, [CloudNativePG](#) using [pgpool-II](#) routing, passive/active and read/write nodes.
- Defined engineering & hiring standards by conducting 40+ technical interviews.

Senior Software Engineer at ethon Zurich, Switzerland

June 2024 - November 2024

MLOps and data on computer vision pipelines. Member of the visual quality inspection team.

- Designed and led the implementation of a distributed system for exchanging vision models and anomaly detection results between training servers and edge devices using [MQTT](#) and different NFSs (e.g. [RustFS](#)) for on-premise and cloud deployments.
- Extended customer-facing UI vision workflows (Angular) to support ROI monitoring, defect heatmaps, annotations, and model evaluation, improving usability and transparency of ML predictions.
- Redesigned vision data storage architecture and reduced storage footprint by 40%.
- Drove improvements to testing infrastructure: UI unit testing, revised integration test framework.

Python

Go

Timescale

Flask

TensorFlow

ONNX

Hydra

CUDA

Docker

[kind](#)

[Gradio](#)

GCP

[Buildkite](#)

Argo

Senior Software Engineer at 1PlusX/TripleLift Zurich, Switzerland

May 2020 - April 2024

Data management and analytics platform providing multi-dimensional batch and real-time user tracking and profiling. Data consolidation and Audience teams.

- Led multi-sprint/multi-quarter and cross-team data projects: Mentoring small groups of developers while resolving ambiguity with non-engineering dependencies such as legal, product, customer success, and external clients.
- Real-time data ingestion ECS infrastructure able to accommodate 7.5M unique events per second on a 30-node Kafka cluster.
- Owned and evolved 50+ unique data pipelines running on production systems with 2PB+ of daily traffic (spawning 30-40k Spark jobs daily). Orchestrated using [Luigi](#). Running on EMR on EKS (or EC2).
- Designed and maintained Terraform-based infrastructure for hundreds of AWS services.
- Tag manager - [Play](#) - ingestion restructuring resulting in a dataset size decrease of up to 80% - while maintaining targeting quality.
- Decreased pipeline runtime up to 15% by introducing batch [Akka](#) API endpoints.

- Defined contracts between data engineering and ML teams, reducing iteration cycles and schema drifts.
- Cross-team user ID storage optimization resulting in up to 20% decrease in Couchbase loads. Load on external APIs was decreased up to 40%.
- Data privacy pipelines: Access and deletion requests' handling of multiple protocols (GDPR, CCPA).
- DSL design for ETL ops: Rich syntax for creating and combining attributes.
- E2E user tracking: From cookies and browser-side TypeScript SDKs to graph-based ID management.
- Angular 2 work on analytics dashboards, widgets, data pipeline builder.

Scala Java Go Python Redux ECR RDS s3 CloudWatch Docker Packer Grafana
InfluxDB DynamoDB PostgreSQL Databricks Elasticsearch OpenSearch

Software Engineer at Veezoo AG Zurich, Switzerland March 2019 - May 2020
NLP platform providing analytics through a variety of corporate data source integrations - namely Oracle, MySQL, PostgreSQL, Salesforce, IBM DB2, Hadoop, SAP HANA - using a QA interface.

- Developed and owned data onboarding automation from raw data to [Apache Jena](#): Knowledge graph ontology representations API for (non)-relational data sources.
- Semantic parsing pipelines for natural languages (bag of words to logical forms to SQL generation).
- Extended and maintained the [Duckling](#) language engine written in Haskell.

Scala Play Solr Docker OpenShift

Software Engineer at Locatee AG Zurich, Switzerland April 2017 - March 2019
Analytics platform providing hyper-location, real-time workspace occupancy insights and, predictive APIs through geospatial data scraping (SNMP protocol).

- Led the modularization of the core [Spring](#) Java 8 app monolith. From creating three distinct services - Analytics API, Data processing service, and pollers - to POCs and eventually to global roll-outs.
- Online data distillation pipeline detecting outliers on workspace behavior and identifying team collaboration coefficients.
- Implemented different MySQL read/write optimizations for measurements post-processing.

InfluxDB Tableau Neo4J Kafka Python Jupyter Pandas Docker

Software Engineer at Upstream Systems SA Athens, Greece October 2014 - March 2017
Green-field project re-engineering the core CRM monolith as a fully distributed system. The implementation included a multitude of APIs for data streaming and job scheduling accommodating thousands to millions of transactions daily.

- Developed asynchronous event-driven services on distributed environments using Guava and RabbitMQ.
- Implementation of Elasticsearch user samplers via Groovy scripts and custom treemap algorithms. Stretch testing, benchmarking, and performance tweaking on Elasticsearch operations.
- Performance profiling ([Dropwizard](#)'s Metrics, YourKit, Grafana, and other tools).
- [Spring Boot](#) API versioning (URL & Resource Representation).
- Batch job framework implementation with RabbitMQ and [Quartz](#). Throttling and flow control features included.
- QA Automation framework (Selenium, Spring).
- [Recommendation engine](#) using Mahout and supervised learning regression algorithms (Log Likelihood Similarity for item-based recommender and Pearson correlation similarity for user-based recommender).

Java 8 Spring JPA TestNG PostgreSQL EhCache flyway Hibernate Camel Kafka Spark
Redis Docker Ansible Protobuf

Software Engineer at Accenture Athens, Greece June 2014 - September 2014
Provisioning and mediation over Oracle databases and mobile telecom software.

TOAD SQL Developer WebLogic

Other Works

- Designer and developer of a [B+ indexed RDBMS in C](#).
- Co-creator of a [social network](#) based on node.js and socket.io. The network provided real-time chatting and financial transactions between users.
- Prototypes for content moderation and text search using Flask and FastAPI. [\[1\]](#) [\[2\]](#)

Education

Degrees

September 2013	M.Sc. in Computer Science Department of Computer Science, National and Kapodistrian University of Athens - 7.49/10 or Upper Second-Class Honours
----------------	--

Languages

Greek	Native
English	Certificate of Proficiency in English, University of Cambridge, Grade A Certificate of Proficiency in English, University of Michigan, Honors
German	B1
French	Diplome d'etudes en langue Francaise (DELF)

Conferences

- [MLOps World](#) 2025

References

Available upon request